

Fleeing Paxos by Raft

...

December 1, 2017

Motivation

“

a pair consisting of an integer and a priest, using a lexicographical ordering, where

$$(13, \Gamma\rho\alpha\tilde{\iota}) < (13, \Lambda\iota\nu\sigma\tilde{\iota}) < (15, \Gamma\rho\alpha\tilde{\iota})$$

...

To maintain B2, a ballot's quorum was chosen to contain a $\mu\alpha\delta\zeta\delta\omega\rho\iota\tilde{\iota}\sigma\tau$ of priests. Initially, $\mu\alpha\delta\zeta\delta\omega\rho\iota\tilde{\iota}\sigma\tau$ just meant a simple majority. Later, it was observed that fat

”

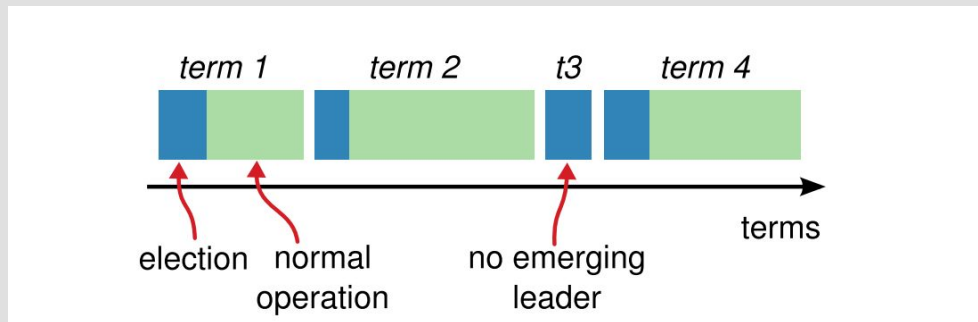
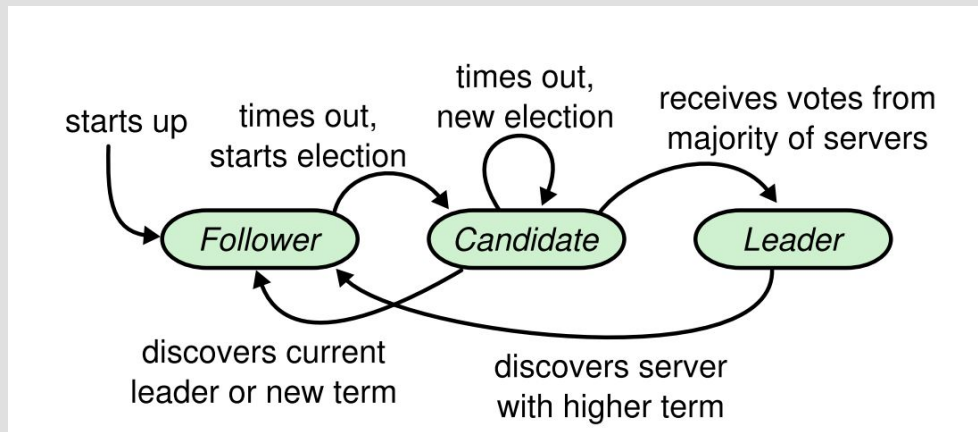
priests were less mobile and spent more time in the Chamber than thin ones,

Raft is Designed for Understandability

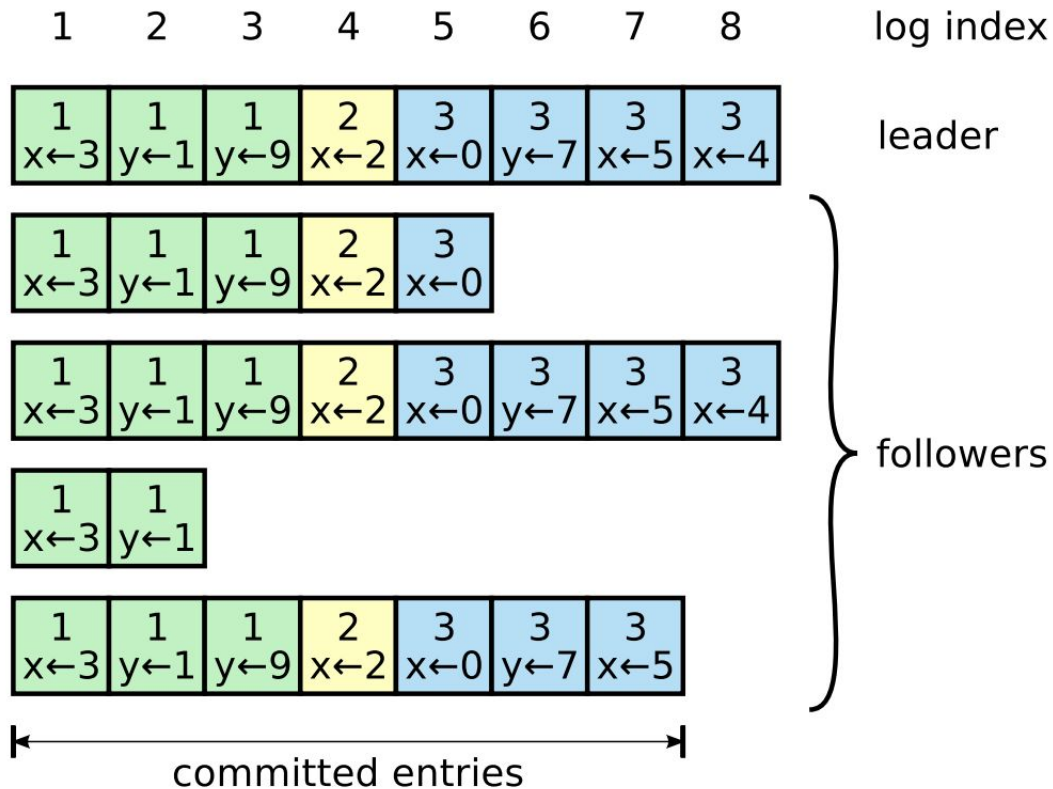
Three largely independent subproblems:

1. Leader Election
2. Log Replication
3. Safety

Structures & States



Structures & States (ctd.)



Subproblems

1. Leader Election

- a. Leader sends heartbeat to followers using empty AppendEntries RPC
- b. Followers each maintain their own randomized countdown timer to declare current leader dead
- c. If timer expires without a heartbeat, Follower increments term, becomes Candidate and sends RequestVote RPC to all other members
- d. If 50% of members reply with their vote, Candidate becomes new Leader

Subproblems

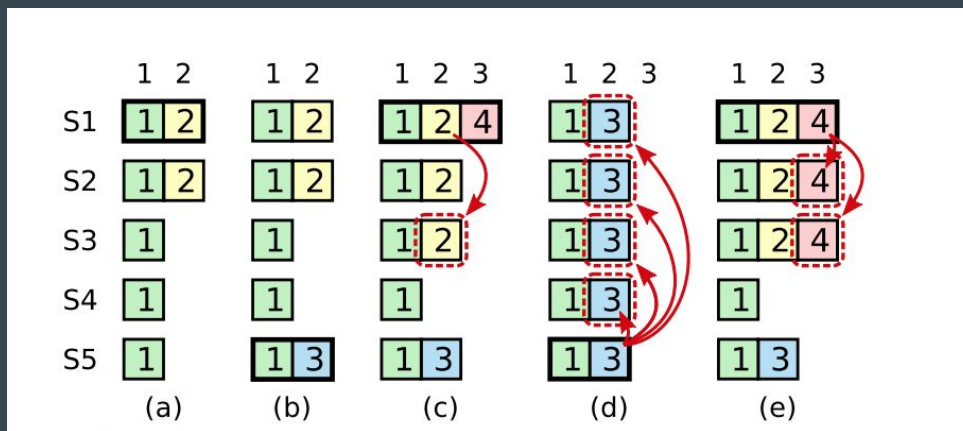
2. Log Replication

- a. All log updates sent by Leader
- b. Same log index && same term value => same log entry
- c. Previous log index and term sent with each update to ensure previous entry is correct
- d. If the previous entry is correct before next entry applied, then all earlier entries are also correct

Subproblems

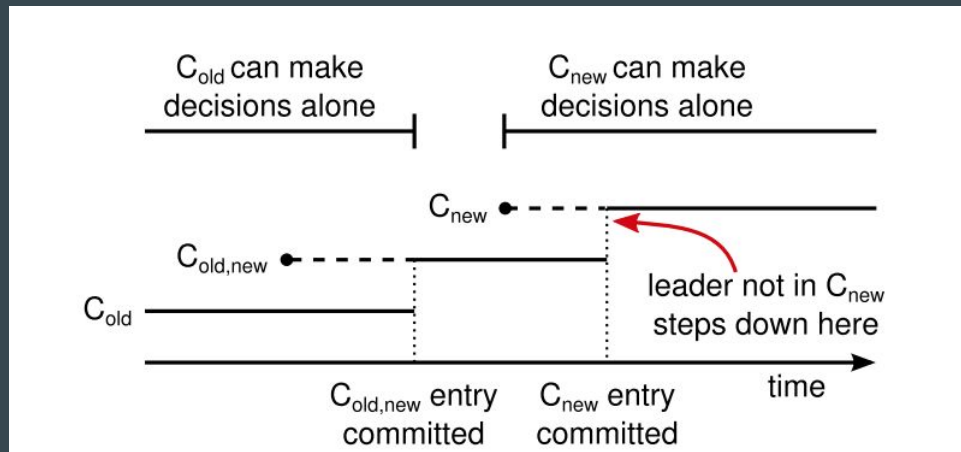
3. Safety

- Only candidates with all log entries committed can be elected leader
- Before granting a vote to a Candidate, the voter checks the Candidate's log for completeness.
- A Candidate with an incomplete log will not be voted in as Leader.
- New leaders will not commit previous term's entries using replica counting



Batteries Included

- Configuration Change
 - Cluster enters joint consensus state
- Log Compaction
- Client Connection
- Linearizability




Demo

Were they successful?

Appendix

Quiz Grades



Survey Results

